

her favorite Internet radio station. An application running on her computer (say, a web browser) has to tell her local router she wants to get frames with this particular multicast address passed her way. The local router must tell the router upstream, and this process continues so each router between the source and destination knows where to pass this multicast data. This ensures that the user can get her rock music without other networks being bothered with this extra data.

PART IV

Transmission Methods

▲CISSP All-in-One Exam Guide

500

IPv4 multicast protocols use a Class D address (224.0.0.0 to 239.255.255.255), which is a special address space reserved for multicasting. IPv6 multicast addresses start

with eight 1's (that is, 1111 1111). Multicasting can be used to send out information;

multimedia data; and even real-time video, music, and voice clips.

Internet Group Management Protocol (IGMP) is used to report multicast group memberships to routers. When a user chooses to accept multicast traffic, she becomes

a member of a particular multicast group. IGMP is the mechanism that allows her computer to inform the local routers that she is part of this group and to send traffic with

a specific multicast address to her system. IGMP can be used for online streaming video

and gaming activities. The protocol allows for efficient use of the necessary resources

when supporting these types of applications.

Like most protocols, IGMP has gone through a few different versions, each improving

upon the earlier one. In version 1, multicast agents periodically send queries to systems on

the network they are responsible for and update their databases, indicating which system

belongs to which group membership. Version 2 provides more granular query types and

allows a system to signal to the agent when it wants to leave a group. Version 3 allows the

systems to specify the sources it wants to receive multicast traffic from. Each version is

backward-compatible because versions 1 and 2 are still in use in legacy equipment.

NOTE The previous statements are true pertaining to IPv4. IPv6 is more than just an upgrade to the original IP protocol; it functions differently in many respects, including how it handles multicasting, which has caused many interoperability issues and delay in its full deployment.

Layer 2 Security Standards

As frames pass from one network device to another device, attackers could sniff the data;

modify the headers; redirect the traffic; spoof traffic; carry out man-in-the-middle attacks,

DoS attacks, and replay attacks; and indulge in other malicious activities. It has become

necessary to secure network traffic at the frame level, which is layer 2 of the OSI model.

802.1AE is the IEEE MAC Security (MACSec) standard, which defines a security infrastructure to provide data confidentiality, data integrity, and data origin authentication.

Where a VPN connection provides protection at the higher networking layers, MACSec

provides hop-by-hop protection at layer 2, as shown in Figure 11-16.

Encrypted

Encrypted

Switch

Encrypted

Switch

Server

Figure 11-16 MACSec provides layer 2 frame protection.

Encrypted

Switch

Workstation

Chapter 11: Networking Fundamentals

501

PART IV

MACSec integrates security protection into wired Ethernet networks to secure LANbased traffic. Only authenticated and trusted devices on the network can communicate

with each other. Unauthorized devices are prevented from communicating via the network, which helps prevent attackers from installing rogue devices and

redirecting

traffic between nodes in an unauthorized manner. When a frame arrives at a device that

is configured with MACSec, the MACSec Security Entity (SecY) decrypts the frame if

necessary and computes an integrity check value (ICV) on the frame and compares it

with the ICV that was sent with the frame. If the ICVs match, the device

processes the frame. If they do not match, the device handles the frame according to a preconfigured policy, such as discarding it.

The IEEE 802.1AR standard specifies unique per-device identifiers (DevID) and the management and cryptographic binding of a device (router, switch, access point) to its identifiers. A verifiable unique device identity allows establishment of the trustworthiness of devices, and thus facilitates secure device provisioning.

As a security administrator you really only want devices that are allowed on your network to be plugged into your network. But how do you properly and uniquely identify devices? The manufacturer's serial number is not available for a protocol to review. MAC addresses, hostnames, and IP addresses are easily spoofed. 802.1AR defines a globally unique per-device secure identifier cryptographically bound to the device through the use of public cryptography and digital certificates. These unique hardwarebased credentials can be used with the Extensible Authentication Protocol-Transport Layer Security (EAP-TLS) authentication framework. Each device that is compliant with IEEE 802.1AR comes with a single built-in initial secure device identity (iDevID). The iDevID is an instance of the general concept of a DevID, which is intended to be used with authentication protocols such as EAP, which is supported by IEEE 802.1X.

So 802.1AR provides a unique ID for a device. 802.1AE provides data encryption, integrity, and origin authentication functionality. 802.1AF carries out key agreement functions for the session keys used for data encryption. Each of these standards provides specific parameters to work within an 802.1X EAP-TLS framework, as shown in Figure 11-17.

As Figure 11-17 shows, when a new device is installed on the network, it cannot just start communicating with other devices, receive an IP address from a Dynamic Host Configuration Protocol (DHCP) server, resolve names with the Domain Name System (DNS) server, and so on. The device cannot carry out any network activity until it is authorized to do so. So 802.1X port authentication kicks in, which means that only authentication data is allowed to travel from the new device to the authenticating server.

The authentication data is the digital certificate and hardware identity associated with that device (802.1AR), which is processed by EAP-TLS. Once the device is authenticated, usually by a Remote Authentication Dial-In User Server (RADIUS) server, encryption

keying material is negotiated and agreed upon between surrounding network devices.

Once the keying material is installed, then data encryption and frame integrity checking

can take place (802.1AE) as traffic goes from one network device to the next.

These IEEE standards are new and evolving and at different levels of implementation

by various vendors. One way the unique hardware identity and cryptographic material

▲CISSP All-in-One Exam Guide

502

IEEE

802.1AF

IEEE

802.1AR

Internal

network

Authentication

server

Certificate

authority

Upstream

device

New

infrastructure

device

IEEE

802.1AE

IETF

Key mgt

framework

0. New device is physically installed.

1. An 802.1X conversation starts.

2. EAP-TLS messages are forwarded.

3. Key material is returned and stored.

4. Session keys are generated.

5. MACSec encryption is enabled.

Figure 11-17 Layer 2 security protocols

are embedded in new network devices is through the use of a Trusted Platform Module

(TPM; described in Chapter 9).

Internet Protocol Networking

Unless your network consists of only a few devices, isolated from the Internet (and what good is that?), you will need to move from layer 2 into layer 3 and above to do anything meaningful. Recall that the data link layer is concerned with exchanging data between devices that are directly connected to each other (in other words, in the same collision domain). Beyond that, we need layer 3 (network) and 4 (transport) protocols, such as TCP/IP. The Transmission Control Protocol/Internet Protocol (TCP/IP) is a suite of protocols that governs the way data travels from one device to another. IP is a network layer protocol and provides datagram routing services. IP's main task is to support internetwork addressing and packet routing. It is a connectionless protocol that envelops data passed to it from the transport layer. The IP protocol addresses the datagram with the source and destination IP addresses. The protocols within the TCP/IP suite work together to break down the data passed from the application layer into pieces that can be moved along a network. They work with other protocols to transmit the data to the destination

▲Chapter 11: Networking Fundamentals

503

IP

IP is a connectionless protocol that provides the addressing and routing capabilities for each package of data. The data, IP, and network relationship can be compared to the relationship between a letter and the postal system:

- Data = Letter
- IP = Addressed envelope
- Network = Postal system

The message is the letter, which is enveloped and addressed by IP, and the network and its services enable the message to be sent from its origin to its destination, like the postal system.

TCP

TCP is referred to as a connection-oriented protocol because before any user data is actually sent, handshaking takes place between the two systems that want to communicate. Once the handshaking completes successfully, a virtual connection is set up between the two systems. UDP is considered a connectionless protocol because it

does not go through these steps. Instead, UDP sends out messages without first contacting the destination computer and does not know if the packets were received properly or dropped. Figure 11-18 shows the difference between a connection-oriented protocol and a connectionless protocol. UDP and TCP sit together on the transport layer, and developers can choose which to use when developing applications. Many times, TCP is the transport protocol of choice because it provides reliability and ensures the packets are delivered. TCP provides a full-duplex, reliable communication mechanism, and if any packets are lost or damaged, they are re-sent; however, TCP requires a lot of system overhead compared to UDP.

PART IV

computer and then reassemble the data back into a form that the application layer can understand and process. Two main protocols work at the transport layer: TCP and UDP. TCP is a reliable and connection-oriented protocol, which means it ensures packets are delivered to the destination computer. If a packet is lost during transmission, TCP has the ability to identify this issue and resend the lost or corrupted packet. TCP also supports packet sequencing (to ensure each and every packet was received), flow and congestion control, and error detection and correction. UDP, on the other hand, is a best-effort and connectionless protocol. It has neither packet sequencing nor flow and congestion control, and the destination does not acknowledge every packet it receives.

▲CISSP All-in-One Exam Guide

504

Connection-oriented communication performs handshaking, sets up a virtual circuit, and verifies that each packet reaches its destination.

Host A

Host B

Connectionless communication just puts the packets on the wire.

Figure 11-18 Connection-oriented protocol vs. connectionless protocol functionality

If developers know that data being dropped during transmission is not detrimental to the application, they may choose to use UDP because it is faster and requires fewer resources. For example, UDP is a better choice than TCP when a server sends status information to all listening nodes on the network. A node will not be negatively affected if, by some chance, it did not receive this status information, because the information will be re-sent every 60 seconds. UDP and TCP are transport protocols that applications use to get their data across a network. They both use ports to communicate with upper OSI layers and to keep track of various conversations that take place simultaneously. The ports are also the mechanism used to identify how other computers access services. When a TCP or UDP message is formed, source and destination ports are contained within the header information along with the source and destination IP addresses. The combination of protocol (TCP or UDP), port, and IP address makes up a socket, and is how packets know where to go (by the address) and how to communicate with the right service or protocol on the other computer (by the port number). The IP address acts as the doorway to a computer, and the port acts as the doorway to the actual protocol or service. To communicate properly, the packet needs to know these doors. Figure 11-19 shows how packets communicate with applications and services through ports.

▲Chapter 11: Networking Fundamentals

505

Web browser

Telnet

SMTP

E-mail client

DNS

HTTP

Ping utility

POP

SSH client

ECHO

FTP

SSH

DNS

SNMP

7

21

22

53

161

Application

23

25

53

80

110

UDP

Transport

Network

Data link

Physical

Figure 11-19 The packet can communicate with upper-layer protocols and services through a port.

The difference between TCP and UDP can also be seen in the message formats. Because TCP offers more services than UDP, it must contain much more information within its packet header format, as shown in Figure 11-20. Table 11-4 lists the major differences between TCP and UDP.

PART IV

TCP

▲CISSP All-in-One Exam Guide

506

Source port

Destination port

Source port

Destination port

Length

Sequence number

Data

Acknowledgment number

Offset Reserved Flags

Window

Checksum

Urgent pointer

Options

Padding

Checksum

UDP format

Data

TCP format

Figure 11-20 TCP carries a lot more information within its segment because it offers more services than UDP.

Property

TCP

UDP

Reliability

Ensures that packets reach their

destinations, returns ACKs when packets are received, and is a reliable protocol.

Does not return ACKs and does not guarantee that a packet will reach its destination. Is an unreliable protocol.

Connection

Connection-oriented. It performs handshaking and develops a virtual connection with the destination computer.

Connectionless. It does no handshaking and does not set up a virtual connection.

Packet sequencing

Uses sequence numbers within headers to make sure each packet within a transmission is received.

Does not use sequence numbers.

Congestion controls

The destination computer can tell the source if it is overwhelmed and thus slow the transmission rate.

The destination computer does not communicate back to the source computer about flow control.

Usage

Used when reliable delivery is required. Intended for relatively small amounts of data transmission.

Used when reliable delivery is not required and high volumes of data need to be transmitted, such as in streaming video and status broadcasts.

Speed and overhead

Uses a considerable amount of

resources and is slower than UDP.

Uses fewer resources and is faster than TCP.

Table 11-4

Major Differences Between TCP and UDP

▲Chapter 11: Networking Fundamentals

507

Port Types

Port numbers up to 1023 (0 to 1023) are called well-known ports, and almost every

computer in the world has the exact same protocol mapped to the exact same port number. That is why they are called well known—everyone follows this same standardized approach. This means that on almost every computer, port 25 is mapped

to SMTP, port 80 is mapped to HTTP, and so on. This mapping between lowernumbered ports and specific protocols is a de facto standard, which just means that

we all do this and that we do not have a standards body dictating that it absolutely

has to be done this way. The fact that almost everyone follows this approach translates to more interoperability among systems all over the world.

Because this is a de facto standard and not a standard that absolutely must be followed, administrators can map different protocols to different port numbers if

that fits their purpose. However, one thing to note is that ports 0 to 1023 can be

used only by privileged system or root processes.

The following shows some of the most commonly used protocols and the ports to which they are usually mapped:

Registered ports are 1024 to 49151, which can be registered with the Internet Assigned Numbers Authority (IANA) for a particular use. Vendors register specific

ports to map to their proprietary software. Dynamic ports (also known as ephemeral

ports) are 49152 to 65535 and are available to be used by any application on an “as needed” basis. Typically, these ports are used on the client side of a connection.

For instance, if you look at a specific connection between your web browser and a

website you visit, you may notice that the destination port is 80 and the source port

(on your client) is 53042. In some cases, however, the server may be listening on a

well-known port (e.g., 135 for RPC) and hand off the server port to an ephemeral one. This means that in some cases, you will see both the client and the server port

for a connection in the range of 49152 to 65535.

PART IV

- Secure Shell (SSH) port 22
- SMTP port 25
- DNS port 53
- HTTP port 80
- NTP port 123
- IMAP port 143
- HTTP Secure (HTTPS) port 443

▲CISSP All-in-One Exam Guide

508

TCP Handshake

TCP must set up a virtual connection between two hosts before any data is sent. This

means the two hosts must agree on certain parameters, data flow, windowing, error

detection, and options. These issues are negotiated during the handshaking phase, as

shown in Figure 11-21.

The host that initiates communication sends a synchronization (SYN) packet to the

receiver. The receiver acknowledges this request by sending a SYN/ACK packet. This packet

translates into, "I have received your request and am ready to communicate with you." The

sending host acknowledges this with an acknowledgment (ACK) packet, which translates

into, "I received your acknowledgment. Let's start transmitting our data." This completes

the handshaking phase, after which a virtual connection is set up, and actual data can

now be passed. The connection that has been set up at this point is considered full duplex,

which means transmission in both directions is possible using the same transmission line.

If an attacker sends a target system SYN packets with a spoofed address, then the victim

system replies to the spoofed address with SYN/ACK packets. Each time the victim system

receives one of these SYN packets, it sets aside resources to manage the new connection.

If the attacker floods the victim system with SYN packets, eventually the victim system

allocates all of its available TCP connection resources and can no longer process new

requests. This is a type of DoS attack that is referred to as a SYN flood. To thwart this type

of attack you can use a number of mitigations, the most common of which are described

in the Internet Engineering Task Force's (IETF) Request for Comments (RFC) 4987. One

of the most effective techniques described in RFC 4987 is the use of SYN caches,

which delays the allocation of a socket until the handshake is completed. Another attack vector we need to understand is TCP sequence numbers. One of the values that is agreed upon during a TCP handshake between two systems is the sequence numbers that will be inserted into the packet headers. Once the sequence number is agreed upon, if a receiving system receives a packet from the sending system that does not have this predetermined value, it disregards the packet. This means that an attacker cannot just spoof the address of a sending system to fool a receiving system; the attacker has to spoof the sender's address and use the correct sequence number values. If an attacker can correctly predict the TCP sequence numbers that two systems will use, then she can create packets containing those numbers and fool the receiving system into thinking that the packets are coming from the authorized sending system. She can then take over the TCP connection between the two systems, which is referred to as TCP session hijacking.

Figure 11-21
The TCP three-way handshake

Host A

1.

SYN

2.

SYN/ACK

3.

ACK

Host B

▲Chapter 11: Networking Fundamentals

509

Data Structures

As stated earlier, the message is formed and passed to the application layer from a program and sent down through the protocol stack. Each protocol at each layer adds its

own information to the message to create a PDU and passes it down to the next layer.

This activity is referred to as encapsulation. As the message is passed down the

stack, it goes through a sort of evolution, and each stage has a specific name that indicates what is taking place. When an application formats data to be transmitted over the network, the PDU is called a message or data. The message is sent to the transport layer, where TCP does its magic on it. The PDU is now a segment. The segment is sent to the network layer. The network layer adds routing and addressing, and now the PDU is called a packet. The network layer passes off the packet to the data link layer, which frames the packet with a header and a trailer, and now it is called a frame. Figure 11-22 illustrates these stages.

EXAM TIP If the message is being transmitted over TCP, it is referred to as a "segment." If it is being transmitted over UDP, it is referred to as a "datagram."

Figure 11-22
Data goes
through its own
evolutionary
stages as it
passes through
the layers within
the network
stack.

Decapsulation

Frame
header

Protocol data units (PDUs)

Data

Application: Data

Transport
header

Data

Transport: Segments

Network
header

Transport
header

Data

Network: Packets

Network
header

Transport
header

Data

Frame
trailer

Data link:
Frames

1 0 1 0 0 1 0 0 0 1 0 1 0 1 0 0 0 1 1 1 0 1 0 1 0 1 0 0 0 0 1 0 0 Bits

PART IV

Sometimes when an author refers to a segment, she is specifying the stage in which the data is located within the protocol stack. If the literature is describing routers, which work at the network layer, the author might use the word “packet” because the data at this layer has routing and addressing information attached. If an author is describing network traffic and flow control, she might use the word “frame” because all data actually ends up in the frame format before it is put on the network wire. The important thing here is that you understand the various steps a data package goes through when it moves up and down the protocol stack.

▲CISSP All-in-One Exam Guide

510

IP Addressing

Each node on a network must have a unique IP address. Today, the most commonly used version of IP is IP version 4 (IPv4), which is used by roughly 70 percent of Internet hosts as we write these words. IP version 6 (IPv6), which was created in part to address the shortage of IPv4 addresses (IPv6 also has many security features built into it that are not part of IPv4), is steadily gaining ground, however. IPv6 is covered later in this chapter. IPv4 uses 32 bits for its addresses, whereas IPv6 uses 128 bits; thus, IPv6 provides

more possible addresses with which to work. Each address has a host portion and a network portion, and the addresses are grouped into classes and then into subnets. The subnet mask of the address differentiates the groups of addresses that define the subnets of a network. IPv4 address classes are listed in Table 11-5.

For any given IP network within an organization, all nodes connected to the network can have different host addresses but a common network address. The host address identifies every individual node, whereas the network address is the identity of the network all the nodes are connected to; therefore, it is the same for each one of them. Any traffic meant for nodes on this network will be sent to the prescribed network address.

A subnet is created from the host portion of an IP address to designate a “sub” network. This allows us to further break the host portion of the address into two or more logical groupings, as shown in Figure 11-23. A network can be logically partitioned to reduce administration headaches, increase traffic performance, and potentially strengthen security. As an analogy, let’s say you work at Toddlers R Us and you are responsible for babysitting 100 toddlers. If you kept all 100 toddlers in one room, you would probably end up crazy. To better manage these kids, you could break them up into groups. The three-year-olds go in the yellow room, the four-year-olds go in the green room, and the five-year-olds go in the blue room. This is what a network administrator would do—break up and separate computer nodes to be able to better control them. Instead of putting them into physical rooms, the administrator puts them into logical rooms (subnets).

To continue with our analogy, when you put your toddlers in different rooms, you would have physical barriers that separate them—walls. Network subnetting is not physical; it is logical. This means you would not have physical walls separating your individual subnets, so how do you keep them separate? This is where subnet masks Class

Address Range

Description

A

0.0.0.0 to 127.255.255.255

The first byte is the network portion, and the remaining 3 bytes are the host portion.

B

128.0.0.0 to 191.255.255.255

The first 2 bytes are the network portion, and the remaining 2 bytes are the host portion.

C

192.0.0.0 to 223.255.255.255

The first 3 bytes are the network portion, and the remaining 1 byte is the host portion.

D

224.0.0.0 to 239.255.255.255

Used for multicast addresses.

E

240.0.0.0 to 255.255.255.255

Reserved for research.

Table 11-5

IPv4 Addressing

▲Chapter 11: Networking Fundamentals

511

Router

Figure 11-23
Subnets create
logical partitions.

Subnet 1
132.201.1.x

Subnet 2
132.201.2.x

Subnet 3
132.201.3.x

PART IV

come into play. A subnet mask defines smaller networks inside a larger network, just like individual rooms are defined within a building. Subnetting allows larger IP address ranges to be divided into smaller, logical,

and more tangible network segments. Consider an organization with several divisions, such as IT, Accounting, HR, and so on. Creating subnets for each division breaks the networks into logical partitions that route traffic directly to recipients without dispersing data all over the network. This drastically reduces the traffic load across the network, reducing the possibility of network congestion and excessive broadcast packets in the network. Implementing network security policies is also much more effective across logically categorized subnets with a demarcated perimeter, as compared to a large, cluttered, and complex network. Subnetting is particularly beneficial in keeping down routing table sizes because external routers can directly send data to the actual network segment without having to worry about the internal architecture of that network and getting the data to individual hosts. This job can be handled by the internal routers, which can determine the individual hosts in a subnetted environment and save the external routers the hassle of analyzing all 32 bits of an IP address and just look at the “masked” bits. TIP You should not have to calculate any subnets for the CISSP exam, but for a better understanding of how this stuff works under the hood, check out the article “IP Tutorial: Subnet Mask and Subnetting” at <https://www.lifewire.com/internet-protocol-tutorial-subnets-818378> (keep in mind that URLs are subject to change from time to time).

▲CISSP All-in-One Exam Guide

512

If the traditional subnet masks are used, they are referred to as classful or classical IP addresses. If an organization needs to create subnets that do not follow these traditional sizes, then it would use classless IP addresses. This just means a different subnet mask would be used to define the network and host portions of the addresses. After it became clear that available IP addresses were running out as more individuals and corporations participated on the Internet, classless interdomain routing (CIDR) was created. A Class B address range is usually too large for most organizations, and a Class C address range is too small, so CIDR provides the flexibility to increase or decrease the class sizes as necessary. CIDR is the method to specify more flexible IP address classes. CIDR is also

referred to as supernetting.

TIP To better understand CIDR, a good resource is “IP Classless Addressing: Classless Inter-Domain Routing (CIDR)/“Supernetting”: www.tcpipguide.com/free/t_IPClasslessAddressingClasslessInterDomainRoutingCI.htm.

Although each node has an IP address, people usually refer to their hostname rather than their IP address. Hostnames, such as www.mheducation.com, are easier for humans to remember than IP addresses, such as 198.105.254.228. However, the use of these two nomenclatures requires mapping between the hostnames and IP addresses because the computer understands only the numbering scheme. This process is addressed in the “Domain Name Service” section later in this chapter.

NOTE IP provides addressing, packet fragmentation, and packet timeouts. To ensure that packets do not continually traverse a network forever, IP provides a Time to Live (TTL) value that is decremented every time the packet passes through a router.

IPv6

IPv6, also called IP next generation (IPng), not only has a larger address space than IPv4

to support more IP addresses; it has some capabilities that IPv4 does not and it accomplishes some of the same tasks differently. All of the specifics of the new functions within

IPv6 are beyond the scope of this book, but we will look at a few of them, because IPv6

is the way of the future. IPv6 allows for scoped addresses, which enables an administrator to restrict specific addresses for specific servers or file and print sharing, for example.

IPv6 has Internet Protocol Security (IPSec) integrated into the protocol stack, which

provides end-to-end secure transmission and authentication. IPv6 has more flexibility

and routing capabilities and allows for Quality of Service (QoS) priority values to be

assigned to time-sensitive transmissions. The protocol offers autoconfiguration, which

makes administration much easier, and it does not require network address translation

(NAT) to extend its address space.

NAT was developed because IPv4 addresses were running out. Although the NAT technology is extremely useful, it has caused a lot of overhead and transmission problems

because it breaks the client/server model that many applications use today. One reason

the industry did not jump on the IPv6 bandwagon when it came out years ago is that

NAT was developed, which reduced the speed at which IP addresses were being depleted.

513

Although the conversion rate from IPv4 to IPv6 is slow in some parts of the world and the implementation process is quite complicated, the industry is making the shift because of all the benefits that IPv6 brings to the table.
NOTE NAT is covered in the “Network Address Translation” section later in this chapter.

The IPv6 specification, as outlined in RFC 8200, lays out the differences and benefits of IPv6 over IPv4. A few of the differences are as follows:

IPv4 header

0

4

8

Version

IHL

12

16

20

Type of service

28

31

Total length

Identification

Time to live

24

Flags

Protocol

Fragment offset

Header checksum

Source address

Destination address

IPv6 header

0

4

8

Version Traffic class

12

16

20

24

Flow label

28

32

36

44

Payload length

Source address

Destination address

Figure 11-24 IPv4 vs. IPv6 headers

40

48

52

56

Next header

60

Hop limit

63

PART IV

- IPv6 increases the IP address size from 32 bits to 128 bits to support more levels of addressing hierarchy, a much greater number of addressable nodes, and simpler

autoconfiguration of addresses.

- The scalability of multicast routing is improved by adding a “scope” field to multicast addresses. Also, a new type of address called an anycast address is defined, which is used to send a packet to any one of a group of nodes.
- Some IPv4 header fields have been dropped or made optional to reduce the common-case processing cost of packet handling and to limit the bandwidth cost of the IPv6 header. This is illustrated in Figure 11-24.

▲CISSP All-in-One Exam Guide

514

- Changes in the way IP header options are encoded allow for more efficient forwarding, less stringent limits on the length of options, and greater flexibility

for introducing new options in the future.

- A new capability is added to enable the labeling of packets belonging to particular traffic “flows” for which the sender requests special handling, such as

nondefault QoS or “real-time” service.

- Extensions to support authentication, data integrity, and (optional) data confidentiality are also specified for IPv6.

IPv4 limits packets to 65,535 bytes of payload, and IPv6 extends this size to 4,294,967,295 bytes. These larger packets are referred to as jumbograms and improve

performance over high-MTU links. Currently most of the world still uses IPv4, but IPv6

is being deployed more rapidly. This means that there are “pockets” of networks using IPv4

and “pockets” of networks using IPv6 that still need to communicate. This communication

takes place through different tunneling techniques, which either encapsulate IPv6

packets within IPv4 packets or carry out automated address translations.

Automatic

tunneling is a technique where the routing infrastructure automatically determines the

tunnel endpoints so that protocol tunneling can take place without preconfiguration.

In the 6to4 tunneling method, the tunnel endpoints are determined by using a wellknown IPv4 anycast address on the remote side and embedding IPv4 address data within

IPv6 addresses on the local side. Teredo is another automatic tunneling technique that

uses UDP encapsulation so that NAT address translations are not affected.

Intra-Site

Automatic Tunnel Addressing Protocol (ISATAP) treats the IPv4 network as a virtual IPv6

local link, with mappings from each IPv4 address to a link-local IPv6 address.

The 6to4 and Teredo are intersite tunneling mechanisms, and ISATAP is an intrasite

mechanism. So the first two are used for connectivity between different networks,

and ISATAP is used for connectivity of systems within a specific network. Notice in

Figure 11-25 that 6to4 and Teredo are used on the Internet and ISATAP is used within an intranet. While many of these automatic tunneling techniques reduce administration overhead, because network administrators do not have to configure each and every system and network device with two different IP addresses, there are security risks that need to be understood. Many times users and network administrators do not know that automatic tunneling capabilities are enabled, and thus they do not ensure that these different tunnels are secured and/or are being monitored. If you are an administrator of a network and have intrusion detection systems (IDSs), intrusion prevention systems (IPSs), and firewalls that are only configured to monitor and restrict IPv4 traffic, then all IPv6 traffic could be traversing your network insecurely. Attackers use these protocol tunnels and misconfigurations to get past these types of security devices so that malicious activities can take place unnoticed. If you are a user and have a host-based firewall that only understands IPv4 and your operating system has a dual IPv4/IPv6 networking stack, traffic could be bypassing your firewall without being monitored and logged. The use of Teredo can actually open ports in NAT devices that allow for unintended traffic in and out of a network.

▲Chapter 11: Networking Fundamentals

515

Client connecting from
globally routable IPv6
address

External CRL
distribution

Domain controllers

NAP servers

Network
location server

DNS servers

Certification

authority

IPv6

6to4

Client connecting from
public IPv4 address

Teredo

Client connecting from
private (NAT) IPv4 address

IPv4

ISATAP tun

neled IPV6

traffic

Internal CRL
distribution point

IPv6

IP-HTTPS

Client connecting from
behind a firewall, or unable
to connect via other
methods

Application servers
running IPv4

Internet

Application servers
running ISATAP

Application servers
running native IPv6

Intranet

It is critical that people who are responsible for configuring and maintaining systems and networks understand the differences between IPv4 and IPv6 and how the various tunneling mechanisms work so that all vulnerabilities are identified and properly addressed. Products and software may need to be updated to address both traffic types, proxies may need to be deployed to manage traffic communication securely, IPv6 should be disabled if not needed, and security appliances need to be configured to

monitor all
traffic types.

Address Resolution Protocol

On a TCP/IP network, each computer and network device requires a unique IP address and a unique physical hardware address. Each NIC has a unique 48-bit physical address that is programmed by the manufacturer into the ROM chips on the card. The physical address is also referred to as the Media Access Control (MAC) address. The network layer works with and understands IP addresses, and the data link layer works with and understands physical MAC addresses. So, how do these two types of addresses work together while operating at different layers?

NOTE A MAC address is unique because the first 24 bits represent the manufacturer code and the last 24 bits represent the unique serial number assigned by the manufacturer.

PART IV

Figure 11-25 Various IPv4 to IPv6 tunneling techniques

▲CISSP All-in-One Exam Guide

516

When data comes from the application layer, it goes to the transport layer for sequence numbers, session establishment, and streaming. The data is then passed to the network layer, where routing information is added to each packet and the source and destination IP addresses are attached to the data bundle. Then this goes to the data link layer, which must find the MAC address and add it to the header portion of the frame. When a frame hits the wire, it only knows what MAC address it is heading toward. At this lower layer of the OSI model, the mechanisms do not even understand IP addresses. So if a computer cannot resolve the IP address passed down from the network layer to the corresponding MAC address, it cannot communicate with that destination computer.

NOTE A frame is data that is fully encapsulated, with all of the necessary headers and trailers.

MAC and IP addresses must be properly mapped so they can be correctly resolved. This happens through the Address Resolution Protocol (ARP). When the data link layer receives a frame, the network layer has already attached the destination IP address to it, but the data link layer cannot understand the IP address and thus invokes ARP for

help. ARP broadcasts a frame requesting the MAC address that corresponds with the destination IP address. Each computer on the broadcast domain receives this frame, and all but the computer that has the requested IP address ignore it. The computer that has the destination IP address responds with its MAC address. Now ARP knows what hardware address corresponds with that specific IP address. The data link layer takes the frame, adds the hardware address to it, and passes it on to the physical layer, which enables the frame to hit the wire and go to the destination computer. ARP maps the hardware address and associated IP address and stores this mapping in its table for a predefined amount of time. This caching is done so that when another frame destined for the same IP address needs to hit the wire, ARP does not need to broadcast its request again. It just looks in its table for this information. Sometimes attackers alter a system's ARP table so it contains incorrect information. This is called ARP table cache poisoning. The attacker's goal is to receive packets intended for another computer. This is a type of masquerading attack. For example, let's say that Bob's computer has an IP address of 10.0.0.1 and a MAC address of bb:bb:bb:bb:bb:bb, Alice's computer has an IP address of 10.0.0.7 and a MAC address of aa:aa:aa:aa:aa:aa, and an attacker has an IP address of 10.0.0.3 and a MAC address of cc:cc:cc:cc:cc:cc, as shown in Figure 11-26. Suppose Bob wants to send a message to Alice. The message is encapsulated at the IP layer with information including Alice's IP address and then handed off to the data link layer. If this is the first message for Alice's computer, the data link process on Bob's computer has no way of knowing her MAC address, so it crafts an ARP query that (literally) says "Who has 10.0.0.7?" This ARP frame is broadcast to the network, where it is received by both Alice's computer and the attacker's computer. Both respond claiming to be the rightful owners of that IP address. What does Bob's computer do when faced with multiple different responses? The answer in most cases is that it uses the most recent response. If the attacker wants to ensure that Bob's ARP table remains poisoned, then he will have to keep pumping out bogus ARP replies.

517
IP: 10.0.0.7
MAC: aa:aa:aa:aa:aa:aa

Alice

Who has
10.0.0.7?
10.0.0.7 is at
aa: aa:aa:aa:aa:aa
Bob

IP: 10.0.0.1
MAC: bb:bb:bb:bb:bb:bb

10.0.0.7 is at
cc:cc:cc:cc:cc:cc:
10.0.0.7 is at
cc:cc:cc:cc:cc:cc:
10.0.0.7 is at
cc:cc:cc:cc:cc:cc:
10.0.0.7 is at
cc:cc:cc:cc:cc:cc:

Attacker

IP: 10.0.0.3
MAC: cc:cc:cc:cc:cc:cc

Figure 11-26 ARP poisoning attack

Dynamic Host Configuration Protocol

A computer can receive its IP addresses in a few different ways when it first boots up. If

it has a statically assigned address, nothing needs to happen. It already has the configuration settings it needs to communicate and work on the intended network. If a computer

depends upon a DHCP server to assign it the correct IP address, it boots up and makes

a request to the DHCP server. The DHCP server assigns the IP address, and everyone is happy.

DHCP is a UDP-based protocol that allows servers to assign IP addresses to network

clients in real time. Unlike static IP addresses, where IP addresses are manually configured,

the DHCP server automatically checks for available IP addresses and correspondingly

assigns an IP address to the client. This eliminates the possibility of IP address conflicts

that occur if two systems are assigned identical IP addresses, which could cause loss of

service. On the whole, DHCP considerably reduces the effort involved in managing large-scale IP networks.

The DHCP server assigns IP addresses in real time from a specified range when a client connects to the network; this is different from static addresses, where each system

PART IV

So ARP is critical for a system to communicate, but it can be manipulated to allow traffic to be sent to unintended systems. ARP is a rudimentary protocol and does not have any security measures built in to protect itself from these types of attacks.

Networks should have IDS sensors monitoring for this type of malicious activity so that

administrators can be alerted if it is underway. This is not difficult to detect, since, as

already noted, the attacker will have to constantly (or at least frequently) transmit bogus

ARP replies.

▲CISSP All-in-One Exam Guide

518

is individually assigned a specific IP address when coming online. In a standard DHCPbased network, the client computer broadcasts a DHCPDISCOVER message on the network in search of the DHCP server. Once the respective DHCP server receives the

DHCPDISCOVER request, the server responds with a DHCPOFFER packet, offering the client an IP address. The server assigns the IP address based on the subject of the

availability of that IP address and in compliance with its network administration policies.

The DHCPOFFER packet that the server responds with contains the assigned IP address

information and configuration settings for client-side services.

Once the client receives the settings sent by the server through the DHCPOFFER packet, it responds to the server with a DHCPREQUEST packet confirming its acceptance of the allotted settings. The server now acknowledges with a DHCPACK packet, which includes the validity period (lease) for the allocated parameters. Client

DHCP server

1. DHCP discover

2. DHCP offer

3. DHCP request

4. DHCP ack

So as shown in Figure 11-27, the DHCP client yells out to the network, “Who can help me get an address?” The DHCP server responds with an offer: “Here is an address and the parameters that go with it.” The client accepts this gracious offer with the

DHCPREQUEST message, and the server acknowledges this message. Now the client can start interacting with other devices on the network and the user can surf the Web and check her e-mail.

Unfortunately, both the client and server segments of DHCP are vulnerable to falsified identity. On the client end, attackers can masquerade their systems to appear as valid network clients. This enables rogue systems to become a part of an organization's network and potentially infiltrate other systems on the network. An attacker may create an unauthorized DHCP server on the network and start responding to clients searching

Chapter 11: Networking Fundamentals

519

Client

DHCP
server

- DHCPDISCOVER—Client searches for the presence of DHCP servers.
 - DHCPOFFER—DHCP server offers the client an available IP address and service settings.
 - DHCPREQUEST—Client confirms accepting the allocated settings.
- Client

DHCP
server

- DHCPACK—DHCP server acknowledges that the particular IP address has now been allocated to the client system for a specific lease period.

Figure 11-27 The four stages of the Discover, Offer, Request, and Acknowledgment (D-O-R-A) process

PART IV

for a DHCP server. A DHCP server controlled by an attacker can compromise client system configurations, carry out man-in-the-middle attacks, route traffic to unauthorized networks, and a lot more, with the end result of jeopardizing the entire network.

An effective method to shield networks from unauthenticated DHCP clients is through the use of DHCP snooping on network switches. DHCP snooping ensures that DHCP servers can assign IP addresses to only selected systems, identified by their MAC addresses. Also, advanced network switches have the capability to direct clients toward legitimate DHCP servers to get IP addresses and restrict rogue systems from

becoming

DHCP servers on the network.

Diskless workstations do not have a full operating system but have just enough code

to know how to boot up and broadcast for an IP address, and they may have a pointer to

the server that holds the operating system. The diskless workstation knows its hardware

address, so it broadcasts this information so that a listening server can assign it the

correct IP address. As with ARP, Reverse Address Resolution Protocol (RARP) frames go to

all systems on the subnet, but only the RARP server responds. Once the RARP server

receives this request, it looks in its table to see which IP address matches the broadcast

hardware address. The server then sends a message that contains its IP address back to

the requesting computer. The system now has an IP address and can function on the

network.

The Bootstrap Protocol (BOOTP) was created after RARP to enhance the functionality

that RARP provides for diskless workstations. The diskless workstation can receive its

IP address, the name server address for future name resolutions, and the default gateway

address from the BOOTP server. BOOTP usually provides more functionality to diskless

workstations than does RARP.

▲CISSP All-in-One Exam Guide

520

Internet Control Message Protocol

The Internet Control Message Protocol (ICMP) is basically IP's "messenger boy."

ICMP

delivers status messages, reports errors, replies to certain requests, and reports routing

information and is used to test connectivity and troubleshoot problems on IP networks.

The most commonly understood use of ICMP is its use by the ping utility. When a person wants to test connectivity to another system, he may ping it, which sends out

ICMP Echo Request frames. The replies on his screen that are returned to the ping utility

are called ICMP Echo Reply frames and are responding to the Echo Request frames.

If a

reply is not returned within a predefined time period, the ping utility sends more Echo

Request frames. If there is still no reply, ping indicates the host is unreachable.

ICMP also indicates when problems occur with a specific route on the network and

tells surrounding routers about better routes to take based on the health and congestion of the various pathways. Routers use ICMP to send messages in response to packets that could not be delivered. The router selects the proper ICMP response and sends it back to the requesting host, indicating that problems were encountered with the transmission request. ICMP is used by other connectionless protocols, not just IP, because connectionless protocols do not have any way of detecting and reacting to transmission errors, as do connection-oriented protocols. In these instances, the connectionless protocol may use ICMP to send error messages back to the sending system to indicate networking problems. As you can see in Table 11-6, ICMP is used for many different networking purposes. This table lists the various messages that can be sent to systems and devices through ICMP.

Attacks Using ICMP

ICMP was developed to send status messages, not to hold or transmit user data. But someone figured out how to insert some data inside of an ICMP packet, which can be used to communicate to an already compromised system. This technique is called ICMP tunneling, and is an older, but still effective, client/server approach that can be used by hackers to set up and maintain covert communication channels to compromised systems. The attacker would target a computer and install the server portion of the tunneling software. This server portion would “listen” on a port, which is the back door an attacker can use to access the system. To gain access and open a remote shell to this computer, an attacker would send commands inside of ICMP packets. This is usually successful because many routers and firewalls are configured to allow ICMP traffic to come and go out of the network, based on the assumption that this is safe because ICMP was developed to not hold any data or a payload. Just as any tool that can be used for good can also be used for evil, attackers commonly use ICMP to redirect traffic. The redirected traffic can go to the attacker’s dedicated system, or it can go into a “black hole.” Routers use ICMP messages to update each other on network link status. An attacker could send a bogus ICMP message with incorrect information, which could cause the routers to divert network traffic to where

the attacker indicates it should go. ICMP is also used as the core protocol for a network tool called Traceroute. Traceroute is used to diagnose network connections, but since it gathers a lot of important network statistics, attackers use the tool to map out a victim's network. This is similar to a burglar

▲Chapter 11: Networking Fundamentals

521

Table 11-6

ICMP Message

Types

Name

0

Echo Reply

1

Unassigned

2

Unassigned

3

Destination Unreachable

4

Source Quench

5

Redirect

6

Alternate Host Address

7

Unassigned

8

Echo Request

9

Router Advertisement

10

Router Solicitation

11

Time Exceeded

12

Parameter Problem

13

Timestamp

14

Timestamp Reply

15

Information Request

16

Information Reply

17

Address Mask Request

18

Address Mask Reply

19

Reserved (for Security)

20–29

Reserved (for Robustness Experiment)

30

Traceroute

31

Datagram Conversion Error

32

Mobile Host Redirect

33

IPv6 Where-Are-You

34

IPv6 I-Am-Here

35

Mobile Registration Request

36

Mobile Registration Reply

37

Domain Name Request

38

Domain Name Reply

39

SKIP

40

Photuris (Disambiguation)

41

ICMP messages utilized by experimental mobility protocols
such as Seamoby

PART IV

Type

♣CISSP All-in-One Exam Guide

522

“casing the joint,” meaning that the more the attacker learns about the environment,
the easier it can be for her to exploit some critical targets. So while the Traceroute tool

is a valid networking program, a security administrator might configure the IDS sensors to monitor for extensive use of this tool because it could indicate that an attacker is attempting to map out the network's architecture. The countermeasures to these types of attacks are to use firewall rules that only allow the necessary ICMP packets into the network and the use of an IDS or IPS to watch for suspicious activities. Host-based protection (host firewalls and host IDS) can also be installed and configured to identify this type of suspicious behavior.

Simple Network Management Protocol

The Simple Network Management Protocol (SNMP) was released to the networking world in 1988 to help with the growing demand of managing network IP devices. Organizations use many types of products that use SNMP to view the status of their network, traffic flows, and the hosts within the network. Since these tasks are commonly carried out using graphical user interface (GUI)-based applications, many people do not have a full understanding of how the protocol actually works. The protocol is important to understand because it can provide a wealth of information to attackers, and you should understand the amount of information that is available to the ones who wish to do you harm, how they actually access this data, and what can be done with it. The two main components within SNMP are managers and agents. The manager is the server portion, which polls different devices to check status information. The server component also receives trap messages from agents and provides a centralized place to hold all network-wide information. The agent is a piece of software that runs on a network device, which is commonly integrated into the operating system. The agent has a list of objects that it is to keep track of, which is held in a database-like structure called the Management Information Base (MIB). A MIB is a logical grouping of managed objects that contain data used for specific management tasks and status checks. When the SNMP manager component polls the individual agent installed on a specific device, the agent pulls the data it has collected from the MIB and sends it to the manager. Figure 11-28 illustrates how data pulled from different devices is located in one centralized location (SNMP manager). This allows the network administrator to have a holistic view of the network and the devices that make up that network. NOTE The trap operation allows the agent to inform the manager of an event, instead of having to wait to be polled. For example, if an interface on

a router goes down, an agent can send a trap message to the manager. This is the only way an agent can communicate with the manager without first being polled.

It might be necessary to restrict which managers can request information of an agent, so communities were developed to establish a trust between specific agents and managers.

A community string is basically a password a manager uses to request data from the agent, and there are two main community strings with different levels of access: read-only and

Chapter 11: Networking Fundamentals

523

SNMP agent

MIB

SNMP agent

MIB

SNMP agent

MIB

SNMP agent

MIB

WAN

SNMP agent

MIB

SNMP agent

MIB

SNMP agent

MIB

Statistics

Alerts

Events

SNMP manager

read-write. As the names imply, the read-only community string allows a manager to read data held within a device's MIB, and the read-write string allows a manager to read the data and modify it. If an attacker can uncover the read-write string, she could change values held within the MIB, which could reconfigure the device. Since the community string is a password, it should be hard to guess and be protected. It should contain mixed-case alphanumeric strings that are not dictionary words.

This practice is not always the case in many networks. The usual default read-only community string is "public" and the read-write string is "private." Many organizations do not change these, so anyone who can connect to port 161 can read the status information of a device and potentially reconfigure it. Different vendors may put in their own default community string values, but organizations may still not take the necessary steps to change them. Attackers usually have lists of default vendor community string values, so they can be easily discovered and used against networks. To make matters worse, the community strings are sent in cleartext in SNMP v1 and v2, so even if a company does the right thing by changing the default values, the strings are still easily accessible to any attacker with a sniffer. If you absolutely have to use v1 or v2 (and you really shouldn't because they are obsolete), make sure that different network segments use different community strings, so that if one string is compromised an attacker cannot gain access to all the devices in the network. The SNMP ports (161 and 162) should not be open to untrusted networks, like the Internet, and if needed they should be filtered to ensure only authorized individuals can connect to them. If these

PART IV

Figure 11-28 Agents provide the manager with SNMP data.

♣CISSP All-in-One Exam Guide

524

ports need to be available to an untrusted network, configure the router or firewall to only allow UDP traffic to come and go from preapproved network-management stations. While versions 1 and 2 of this protocol send the community string values in cleartext, version 3 has cryptographic functionality, which provides encryption, message integrity, and authentication security. So, SNMP v3 should be implemented for more granular protection. If the proper countermeasures are not put into place, then an attacker can gain access to a wealth of device-oriented data that can be used in her follow-up attacks. The following are just some data sets held within MIB SNMP objects that attackers would be interested in:

- .server.svSvcTable.svSvcEntry.svSvcName

Running services

.server.svShareTable.svShareEntry.svShareName

Share names

.server.sv.ShareTable.svShareEntry.svSharePath

Share paths

.server.sv.ShareTable.svShareEntry.svShareComment

Comments on shares

.server.svUserTable.svUserEntry.svUserName

Usernames

.domain.domPrimaryDomain8

Domain names

Gathering this type of data allows an attacker to map out the target network and enumerate the nodes that make up the network.

As with all tools, SNMP is used for good purposes (network management) and for bad purposes (target mapping, device reconfiguration). We need to understand both sides of all tools available to us.

Domain Name Service

Imagine how hard it would be to use the Internet if we had to remember actual specific

IP addresses to get to various websites. The Domain Name Service (DNS) is a method of

resolving hostnames to IP addresses so names can be used instead of IP addresses within

networked environments.

The first iteration of the Internet was made up of about 100 computers (versus over 22 billion now), and a list was kept that mapped every system's hostname to its

IP address. This list was kept on an FTP server so everyone could access it. It did not take

long for the task of maintaining this list to become overwhelming, and the computing

community looked to automate it.

When a user types a uniform resource locator (URL) into his web browser, the URL is made up of words or letters that are in a sequence that makes sense to that user, such

as www.google.com. However, these words are only for humans—computers work with IP addresses. So after the user enters this URL and presses `<enter>`, behind the scenes

his computer is actually being directed to a DNS server that will resolve this

URL, or
hostname, into an IP address that the computer understands. Once the hostname
has
been resolved into an IP address, the computer knows how to get to the web
server
holding the requested web page.

▲Chapter 11: Networking Fundamentals

525

PART IV

Many organizations have their own DNS servers to resolve their internal
hostnames.
These organizations usually also use the DNS servers at their Internet service
providers
(ISPs) to resolve hostnames on the Internet. An internal DNS server can be used
to
resolve hostnames on the entire LAN, but usually more than one DNS server is
used so
the load can be split up and so redundancy and fault tolerance are in place.
Within DNS servers, DNS namespaces are split up administratively into zones.
One zone may contain all hostnames for the marketing and accounting departments,
and another zone may contain hostnames for the administration, research, and
legal
departments. The DNS server that holds the files for one of these zones is said
to be
the authoritative name server for that particular zone. A zone may contain one
or more
domains, and the DNS server holding those host records is the authoritative name
server
for those domains.
The DNS server contains records that map hostnames to IP addresses, which are
referred to as resource records. When a user's computer needs to resolve a
hostname to an
IP address, it looks to its networking settings to find its DNS server. The
computer then
sends a request, containing the hostname, to the DNS server for resolution. The
DNS
server looks at its resource records and finds the record with this particular
hostname,
retrieves the address, and replies to the computer with the corresponding IP
address.
It is recommended that a primary and a secondary DNS server cover each zone. The
primary DNS server contains the actual resource records for a zone, and the
secondary
DNS server contains copies of those records. Users can use the secondary DNS
server
to resolve names, which takes a load off of the primary server. If the primary
server goes
down for any reason or is taken offline, users can still use the secondary
server for name
resolution. Having both a primary DNS server and a secondary DNS server provides

fault

tolerance and redundancy to ensure users can continue to work if something happens to one of these servers.

The primary and secondary DNS servers synchronize their information through a zone transfer. After changes take place to the primary DNS server, those changes must

be replicated to the secondary DNS server. It is important to configure the DNS server

to allow zone transfers to take place only between the specific servers. For years now,

attackers have been carrying out unauthorized zone transfers to gather very useful

network information from victims' DNS servers.

An unauthorized zone transfer provides the attacker with information on almost every

system within the network, including the hostname and IP address of each system, system

alias names, public key infrastructure (PKI) server, DHCP server, DNS servers, and so

on. This allows an attacker to carry out very targeted attacks on specific systems. If you

were the attacker and you had a new exploit for DHCP software, now you would know

the IP address of the company's DHCP server and could send your attack parameters

directly to that system. Also, since the zone transfer can provide data on all of the systems

in the network, the attacker can map out the network. He knows what subnets are being used, which systems are in each subnet, and where the critical network systems

reside. This is analogous to you allowing a burglar into your house with the freedom of

identifying where you keep your jewels, expensive stereo equipment, piggy bank, and

keys to your car, which will allow him to more easily steal these items when you are on

▲CISSP All-in-One Exam Guide

526

vacation. Unauthorized zone transfers can take place if the DNS servers are not properly configured to restrict this type of activity.

Internet DNS and Domains

Networks on the Internet are connected in a hierarchical structure, as are the different

DNS servers, as shown in Figure 11-29. While performing routing tasks, if a router does

not know the necessary path to the requested destination, that router passes the packet

up to a router above it. The router above it knows about all the routers below it. This

router has a broader view of the routing that takes place on the Internet and has a better chance of getting the packet to the correct destination. This holds true with DNS servers also. If one DNS server does not know which DNS server holds the necessary record to resolve a hostname, it can pass the request up to a DNS server above it.

Root-level
domain

Top-level domain

COM

EDU

ORG

Second-level domain

mheducation

usma

ieee

Figure 11-29 The DNS naming hierarchy is similar to the routing hierarchy on the Internet.

▲Chapter 11: Networking Fundamentals

527

The naming scheme of the Internet resembles an inverted tree with the root servers at the top. Lower branches of this tree are divided into top-level domains, with second-level domains under each. The most common top-level domains are as follows:

- COM Commercial
- EDU Education
- MIL U.S. military organization
- INT International treaty organization
- GOV Government
- ORG Organizational
- NET Networks

DNS Resolution Components

Your computer has a DNS resolver, which is responsible for sending out requests to DNS servers for host IP address information. If your system did not have this resolver, when you type in `www.google.com` in your browser, you would not get to this website because your system does not actually know what `www.google.com` means. When you type in this URL, your system's resolver has the IP address of

a DNS server it is supposed to send its hostname-to-IP address request to. Your resolver can send out a nonrecursive query or a recursive query to the DNS server.

A nonrecursive query means that the request just goes to that specified DNS server

and either the answer is returned to the resolver or an error is returned. A recursive

query means that the request can be passed on from one DNS server to another one until the DNS server with the correct information is identified.

(Continued)

PART IV

So how do all of these DNS servers play together in the Internet playground?

When

a user types in a URL to access a website that sells computer books, for example, his

computer asks its local DNS server if it can resolve this hostname to an IP address. If

the primary DNS server cannot resolve the hostname, it must query a higher-level DNS

server, ultimately ending at an authoritative DNS server for the specified domain. Because

this website is most likely not on the corporate network, the local LAN DNS server will

not usually know the necessary IP address of that website. The DNS server does not

reject the user's request, but rather passes it on to another DNS server on the Internet.

The request for this hostname resolution continues through different DNS servers until

it reaches one that knows the IP address. The requested host's IP information is reported

back to the user's computer. The user's computer then attempts to access the website

using the IP address, and soon the user is buying computer books, happy as a clam.

DNS server and hostname resolution is extremely important in corporate networking

and Internet use. Without it, users would have to remember and type in the IP address

for each website and individual system instead of the name. That would be a mess.

▲CISSP All-in-One Exam Guide

528

In the following illustration, you can follow the succession of requests that commonly takes place. Your system's resolver first checks to see if it already has the

necessary hostname-to-IP address mapping cached or if it is in a local HOSTS file.

If the necessary information is not found, the resolver sends the request to the local

DNS server. If the local DNS server does not have the information, it sends the request to a different DNS server.

DNS client (resolver)

Server-to-server
query (recursion)

Client-to-server query
Zones

Other DNS servers
Q3

Q1

DNS
resolver
cache

A1

A3

DNS
server
Q2

Q5
A5

A2
Q4

Web browser
A4

URL: www.google.com

HOSTS file

DNS server
cache

The HOSTS file resides on the local computer and can contain static hostnameto-IP address mapping information. If you do not want your system to query a DNS server, you can add the necessary data in the HOSTS file, and your system will check its contents before reaching out to a DNS server. HOSTS files are like two-edged swords: on the one hand they offer a degree of security by ensuring that certain hosts resolve to specific IP addresses, but on the other hand they are attractive targets for attackers who want to redirect your traffic to specific hosts. The

key, as always, is to carefully analyze and mitigate the risks.

Chapter 11: Networking Fundamentals

529

DNS Threats

PART IV

As stated earlier, not every DNS server knows the IP address of every hostname it is asked to resolve. When a request for a hostname-to-IP address mapping arrives at a DNS server (server A), the server reviews its resource records to see if it has the necessary information to fulfill this request. If the server does not have a resource record for this hostname, it forwards the request to another DNS server (server B), which in turn reviews its resource records and, if it has the mapping information, sends the information back to server A. Server A caches this hostname-to-IP address mapping in its memory (in case another client requests it) and sends the information on to the requesting client. With the preceding information in mind, consider a sample scenario. Andy the attacker wants to make sure that any time one of his competitor's customers tries to visit the competitor's website, the customer is instead pointed to Andy's website. Therefore, Andy installs a tool that listens for requests that leave DNS server A asking other DNS servers if they know how to map the competitor's hostname to its IP address. Once Andy sees that server A sends out a request to server B to resolve the competitor's hostname, Andy quickly sends a message to server A indicating that the competitor's hostname resolves to Andy's website's IP address. Server A's software accepts the first response it gets, so server A caches this incorrect mapping information and sends it on to the requesting client. Now when the client tries to reach Andy's competitor's website, she is instead pointed to Andy's website. This will happen subsequently to any user who uses server A to resolve the competitor's hostname to an IP address because this information is cached on server A. Previous vulnerabilities that have allowed this type of activity to take place have been addressed, but this type of attack is still taking place because when server A receives a response to its request, it does not authenticate the sender. Mitigating DNS threats consists of numerous measures, the most important of

which is the use of stronger authentication mechanisms such as the DNSSEC (DNS security, which is part of many current implementations of DNS server software). DNSSEC implements PKI and digital signatures, which allows DNS servers to validate the origin of a message to ensure that it is not spoofed and potentially malicious. If DNSSEC were enabled on server A, then server A would, upon receiving a response, validate the digital signature on the message before accepting the information to make sure that the response is from an authorized DNS server. So even if an attacker sends a message to a DNS server, the DNS server would discard it because the message would not contain a valid digital signature. DNSSEC allows DNS servers to send and receive authorized messages between themselves and thwarts the attacker's goal of poisoning a DNS cache table. This sounds simple enough, but for DNSSEC to be rolled out properly, all of the DNS servers on the Internet would have to participate in a PKI to be able to validate digital signatures. The implementation of Internet-wide PKIs simultaneously and seamlessly has proved to be difficult. Despite the fact that DNSSEC requires more resources than the traditional DNS, more and more organizations globally are opting to use DNSSEC. As of this writing, 91 percent of the top-level domains implement DNSSEC. However, across the entire Internet, barely 3 percent of domains have implemented it. So we are getting there, slowly but surely.

▲CISSP All-in-One Exam Guide

530

DNS Splitting

Organizations should implement split DNS, which means a DNS server in the DMZ handles external hostname-to-IP address resolution requests, while an internal DNS server handles only internal requests. This helps ensure that the internal

DNS server has layers of protection and is not exposed by being "Internet facing."

The internal DNS server should only contain resource records for the internal computer systems, and the external DNS server should only contain resource records

for the systems the organization wants the outside world to be able to connect to.

If the external DNS server is compromised and it has the resource records for all of the internal systems, now the attacker has a lot of "inside knowledge" and can carry

out targeted attacks. External DNS servers should only contain information on the systems within the DMZ that the organization wants others on the Internet to be able to communicate with (web servers, external mail server, etc.).

Now let's discuss another (indirectly related) predicament in securing DNS traffic—manipulation of the HOSTS file, a technique frequently used by malware. The HOSTS file is used by the operating system to map hostnames to IP addresses as described before.

The HOSTS file is a plaintext file located in the %systemroot%\system32\drivers\etc folder in Windows, in /etc/hosts in Unix/Linux systems, and in /private/etc/hosts in macOS. The HOSTS file simply consists of a list of IP addresses with their corresponding hostnames.

Depending on its configuration, the computer refers to the HOSTS file before issuing a DNS request to a DNS server. Most operating systems give preference to details of IP addresses returned by the HOSTS file rather than querying the DNS server because the HOSTS file is generally under the direct control of the local system administrator.

As covered previously, in the early days of the Internet and prior to the adoption of DNS, HOSTS files were the primary source of determining a host's network addresses from its hostname. With the increase in the number of hosts connected to the Internet, maintaining HOSTS files became next to impossible and ultimately led to the creation of DNS.

Due to the important role of HOSTS files, they are frequently targeted by malware to propagate across systems connected on a local network. Once a malicious program takes over the HOSTS file, it can divert traffic from its intended destination to websites hosting malicious content, for example. A common example of HOSTS file manipulation carried out by malware involves blocking users from visiting antivirus update websites.

This is usually done by mapping target hostnames to the loopback interface IP address 127.0.0.1. The most effective technique for preventing HOSTS file intrusions is to set it as a read-only file and implement a host-based IDS that watches for critical file modification attempts.

Attackers don't always have to go through all this trouble to divert traffic to rogue destinations. They can also use some very simple techniques that are

surprisingly effective
in routing naive users to unintended destinations. The most common approach is known

▲Chapter 11: Networking Fundamentals

531

as URL hiding. Hypertext Markup Language (HTML) documents and e-mail messages allow users to attach or embed hyperlinks in any given text, such as the “Click Here”

links you commonly see in e-mail messages or web pages. Attackers misuse hyperlinks to

deceive unsuspecting users into clicking rogue links.

Let’s say a malicious attacker creates an unsuspecting text, `www.good.site`, but embeds

the link to an abusive website, `www.bad.site`. People are likely to click the `www.good`

`.site` link without knowing that they are actually being taken to the bad site.

In addition,

attackers also use character encoding to obscure web addresses that may arouse user

suspicion.

Network Address Translation

- `10.0.0.0–10.255.255.255` Class A networks
- `172.16.0.0–172.31.255.255` Class B networks
- `192.168.0.0–192.168.255.255` Class C networks

NAT is a gateway that lies between a network and the Internet (or another network)

that performs transparent routing and address translation. Because IP addresses were

depleting fast, IPv6 was developed in 1999, and was intended to be the long-term fix to

the address shortage problem. NAT was developed as the short-term fix to enable more

organizations to participate on the Internet. However, to date, IPv6 is slow in acceptance

and implementation, while NAT has caught on like wildfire. Many firewall vendors have

implemented NAT into their products, and it has been found that NAT actually provides

a great security benefit. When attackers want to hack a network, they first do what they

can to learn all about the network and its topology, services, and addresses. Attackers

cannot easily find out an organization’s address scheme and its topology when NAT is

in place, because NAT acts like a large nightclub bouncer by standing in front of the

network and hiding the true IP scheme.

PART IV

When computers need to communicate with each other, they must use the same type of addressing scheme so everyone understands how to find and talk to one another. The Internet uses the IP address scheme as discussed earlier in the chapter, and any computer or network that wants to communicate with other users on the network must conform to this scheme; otherwise, that computer will sit in a virtual room with only itself to talk to. However, IP addresses have become scarce (until the full adoption of IPv6) and expensive. So some smart people came up with network address translation (NAT), which enables a network that does not follow the Internet's addressing scheme to communicate over the Internet. Private IP addresses have been reserved for internal LAN address use, as outlined in RFC 1918. These addresses can be used within the boundaries of an organization, but they cannot be used on the Internet because they will not be properly routed. NAT enables an organization to use these private addresses and still be able to communicate transparently with computers on the Internet. The following lists current private IP address ranges:

▲CISSP All-in-One Exam Guide

532

NAT hides internal addresses by centralizing them on one device, and any frames that leave that network have only the source address of that device, not of the actual internal computer that sends the message. So when a message comes from an internal computer with the address of 10.10.10.2, for example, the message is stopped at the device running NAT software, which happens to have the IP address of 1.2.3.4. NAT changes the header of the packet from the internal address, 10.10.10.2, to the IP address of the NAT device, 1.2.3.4. When a computer on the Internet replies to this message, it replies to the address 1.2.3.4. The NAT device changes the header on this reply message to 10.10.10.2 and puts it on the wire for the internal user to receive. Three basic types of NAT implementations can be used:

- **Static mapping** The NAT software has a pool of public IP addresses configured. Each private address is statically mapped to a specific public address. So computer A always receives the public address x, computer B always receives the public address y, and so on. This is generally used for servers that need to

keep the same public address at all times.

- Dynamic mapping The NAT software has a pool of IP addresses, but instead of statically mapping a public address to a specific private address, it works on a

first-come, first-served basis. So if Bob needs to communicate over the Internet,

his system makes a request to the NAT server. The NAT server takes the first IP address on the list and maps it to Bob's private address. The balancing act is to estimate how many computers will most likely need to communicate outside the internal network at one time. This estimate is the number of public addresses

the organization purchases, instead of purchasing one public address for each computer.

- Port address translation (PAT) The organization owns and uses only one public IP address for all systems that need to communicate outside the internal network. How in the world could all computers use the exact same IP address? Good question. Here's an example: The NAT device has an IP address of 127.50.41.3. When computer A needs to communicate with a system on the Internet, the NAT device documents this computer's private address and source port number (10.10.44.3; port 43,887). The NAT device changes the IP address in the computer's packet header to 127.50.41.3, with the source port 40,000. When computer B also needs to communicate with a system on the Internet, the NAT device documents the private address and source port number (10.10.44.15; port 23,398) and changes the header information to 127.50.41.3 with source port 40,001. So when a system responds to computer A, the packet first goes to the NAT device, which looks up the port number 40,000 and sees that it maps to computer A's real information. So the NAT device changes the header information to address 10.10.44.3 and port 43,887 and sends it to computer A for processing. An organization can save a lot more money by using PAT because it needs to buy only a few public IP addresses, which are used by all systems in the network.

▲Chapter 11: Networking Fundamentals

533

Most NAT implementations are stateful, meaning they keep track of a communication

between the internal host and an external host until that session is ended. The NAT

device needs to remember the internal IP address and port to send the reply messages

back. This stateful characteristic is similar to stateful-inspection firewalls, but NAT does

not perform scans on the incoming packets to look for malicious characteristics. Instead,

NAT is a service usually performed on routers or gateway devices within an organization's screened subnet.

Although NAT was developed to provide a quick fix for the depleting IP address problem, it has actually put the problem off for quite some time. The more organizations

that implement private address schemes, the less likely IP addresses will become scarce.

This has been helpful to NAT and the vendors that implement this technology, but

it has
put the acceptance and implementation of IPv6 much farther down the road.

Routing Protocols

NOTE As an analogy, just as the world is made up of different countries, the Internet is made up of different ASs. Each AS has delineated boundaries just as countries do. Countries can have their own languages (e.g., Spanish, Arabic, Russian). Similarly, ASs have their own internal routing protocols. Countries that speak different languages need to have a way of communicating with each other, which could happen through interpreters. ASs need to have a standardized method of communicating and working together, which is where external routing protocols come into play.

The architecture of the Internet that supports these various ASs is created so that no entity that needs to connect to a specific AS has to know or understand the interior routing protocols that are being used. Instead, for ASs to communicate, they just have to be using the same exterior routing protocols (see Figure 11-30). As an analogy, suppose you want to deliver a package to a friend who lives in another state. You give the package to your brother, who is going to take a train to the edge of the state and hand it to the postal system at that junction. Thus, you know how your brother will arrive at the edge of the state—by train. You do not know how the postal system will then deliver your package to your friend's house (truck, car, bus), but that is not your concern. It will get to its destination without your participation. Similarly, when one network communicates with another network, the first network puts the data packet (package) on an exterior protocol (train), and when the data packet gets to the border router (edge of the state), the data is transferred to whatever interior protocol is being used on the receiving network.

PART IV

Individual networks on the Internet are referred to as autonomous systems (ASs). These ASs are independently controlled by different service providers and organizations. An AS is made up of routers, which are administered by a single entity and use a common Interior Gateway Protocol (IGP) within the boundaries of the AS. The boundaries of these ASs are delineated by border routers. These routers connect to the border routers of other ASs and run interior and exterior routing protocols. Internal routers